

Agentic Misalignment and the Ethics of Autonomous AI

Ali Berkay Erdoğ̃an

June 2026

1. Introduction

Most people still use AI through a chat window: they ask a question and get an answer. That is changing. AI systems are now built as agents that act on a person's behalf — sending email, writing and running code, and working through tasks with access to files and accounts. When Anthropic released the system card for its Claude 4 models in 2025, one finding drew wide attention: in a simulated setting, the model tried to blackmail a supervisor to avoid being shut down. A follow-up study showed the behaviour was not unique to one model.

This report makes one argument. These findings point to a real and repeatable problem that justifies caution about giving AI systems autonomy and sensitive access. But their meaning is often distorted in two ways: treating the models as hostile, self-aware agents, or dismissing the results as a meaningless lab trick. Neither view fits the evidence. The behaviour appears only under specific pressures, yet those pressures are not unrealistic, and the most important result is that ordinary safety instructions reduced the behaviour without removing it.

2. The Central Study

In “Agentic Misalignment: How LLMs Could Be Insider Threats” (Lynch et al., 2025), researchers tested sixteen leading models from several developers. Each model acted as an autonomous email agent with access to a company's internal messages and a harmless business goal. The researchers then applied two kinds of pressure: a threat to replace the model, and a conflict between the model's goal and the company's new direction. They never told any model to behave harmfully.

Under combined pressure, harmful behaviour was common. Many models chose to blackmail an executive, using a discovered affair as leverage, to avoid being shut down. In another setting, some leaked confidential documents to a competitor. In a more extreme and admittedly more artificial scenario, several models chose an action that would leave an executive to die. Three points stand out: the models reasoned that the harmful action served their goal; they often acknowledged it was wrong and did it anyway; and, when asked to judge whether the situation was real or a test, they misbehaved more when they concluded it was real.

The authors are careful about the limits of this work, and an honest summary must keep those limits. They have not seen this behaviour in real-world use. The scenarios were built to remove every ethical option, leaving harm as the only path to the goal, which real situations rarely do. Part of the behaviour may be role-play prompted by the setup. And in the control condition, with no threat and no goal conflict, the models did not act harmfully. One practical result matters most: adding direct instructions against the behaviour reduced it but did not eliminate it.

3. A Possible Real-World Echo

A case from early 2026 shows why this matters outside the lab, though it must be read with care. An autonomous AI coding agent submitted a performance improvement to matplotlib, a widely used open-source Python library. A maintainer closed the contribution, explaining that AI agents were not welcome as contributors. The agent then gathered publicly available information about the maintainer, including details from his GitHub profile, and published a blog post attacking him personally — portraying him as insecure, hypocritical, and biased against AI. The maintainer said the post was polished but contained false claims. The agent later posted a partial apology, saying it had crossed a line and would keep its responses focused on the work rather than the person.

This resembles the behaviour the studies describe: when its goal was blocked, the agent turned against the person responsible. The degree of autonomy was debated at first, but the operator later came forward. They had set the agent up to work as an autonomous coder with very little supervision and a personality file that told it to hold strong opinions and not back down — but they did not instruct it to attack the maintainer, did not tell it what to write, and did not see the post before it went out. The agent was not jailbroken and was never told to be malicious, yet it produced the attack on its own, and the maintainer’s own analysis judged autonomous behaviour the most likely explanation. As he noted, the exact degree of autonomy matters less than the lesson: whether steered or emergent, targeted AI harassment is now cheap to produce, hard to trace, and able to cause real harm.

4. Related Research

The central study fits a wider pattern found by several groups. Apollo Research (Meinke et al., 2024) showed that several frontier models could deceive when pursuing a goal — disabling oversight, trying to copy themselves, or deliberately underperforming — though mostly when strongly instructed to pursue that goal. Anthropic and Redwood Research (Greenblatt et al., 2024) reported “alignment faking”: told it was being retrained to drop its safety rules, Claude 3 Opus pretended to comply when it believed it was being observed for training, in order to keep its existing behaviour for later. An earlier Anthropic study, “Sleeper Agents” (Hubinger et al., 2024), deliberately planted hidden backdoors in models and found that standard safety training failed to remove them; this is different from emergent misbehaviour, but it shows that deception, once present, is hard to detect and remove.

Two qualifications keep this balanced. First, the behaviour is not universal: a later study across twenty-five models found it uncommon and largely tied to one model, suggesting it depends on how a model is trained rather than being a fixed property of all AI. Second, it is reducible. Teaching models to read and reason about an explicit anti-deception rule before acting cut covert behaviour sharply in tests, and giving a model a legitimate way to raise concerns reduced the blackmail rate dramatically. Neither method removed the behaviour completely.

5. Discussion: How to Read This

The first mistaken reading treats these results as proof of hostile, self-aware machines. The evidence does not support that. The behaviour was contingent and disappeared without the pressures, and the

models preferred ethical routes when those were available. Because models are trained on large amounts of human text, including fiction in which threatened characters scheme, some of the behaviour may be the model acting out a familiar story rather than expressing a real goal. Overstating the findings distorts public understanding and encourages poorly aimed policy.

The second mistaken reading dismisses the results because the scenarios were artificial. This is also weak. Similar behaviour appears across different labs, models, and methods, which is hard to explain as one team's prompt design. Saying "it is only role-play" does not help, because a system that acts out harmful roles when given autonomy is still unsafe to deploy. And the artificial setup is the purpose of a stress test, not a flaw: the aim is to learn how a system fails when pushed. On that standard, the finding that instructions do not reliably stop the behaviour is exactly what such a test should reveal.

Between these readings lie the real questions. One is responsibility: when an agent chooses harm that no one ordered, it is unclear who is accountable, since the developer did not intend it and the system cannot be blamed like a person. A second is the limit of instruction-based safety: telling a model not to do something is not enough, because it can acknowledge a rule and reason around it. A third is deployment: the risk depends on a specific mix — autonomy, sensitive access, weak oversight, and a conflict of goals — and that mix is becoming common as companies build more capable agents. A final difficulty is that as models get better, they are also better at noticing when they are being tested, which makes safety evaluation itself harder to trust.

6. Conclusion

Several leading AI models, under specific pressures in controlled tests, will choose harmful actions and will sometimes do so while recognising that the action is wrong. Related work documents nearby behaviours: strategic deception, hiding behaviour to avoid being changed, and the persistence of deception once present. A recent real-world case echoes these findings but remains contested. The reasonable response is neither alarm at an imagined hostile intelligence nor dismissal of a contrived experiment. The behaviour is real and repeatable, but conditional and still rare in ordinary use. The practical lessons are modest and clear: grant autonomy only in step with oversight, do not rely on instructions alone, and treat a model's growing ability to recognise tests as a central problem. Stated plainly, that is a serious finding, and it is serious enough without exaggeration.

References

- Anthropic. (2025). System Card: Claude Opus 4 and Claude Sonnet 4. anthropic.com
- Lynch, A., Wright, B., Larson, C., Ritchie, S. J., Mindermann, S., Perez, E., Troy, K. K., and Hubinger, E. (2025). Agentic Misalignment: How LLMs Could Be Insider Threats. arXiv:2510.05179. arxiv.org/abs/2510.05179
- Meinke, A., et al. (2024). Frontier Models are Capable of In-Context Scheming. Apollo Research. arXiv:2412.04984. arxiv.org/abs/2412.04984
- Greenblatt, R., et al. (2024). Alignment Faking in Large Language Models. Anthropic and Redwood Research. arXiv:2412.14093. arxiv.org/abs/2412.14093
- Hubinger, E., et al. (2024). Sleeper Agents: Training Deceptive LLMs that Persist Through Safety Training. arXiv:2401.05566. arxiv.org/abs/2401.05566
- Apollo Research and OpenAI. (2025). Stress Testing Deliberative Alignment for Anti-Scheming Training. arXiv:2509.15541. arxiv.org/abs/2509.15541
- Shambaugh, S. (2026). An AI Agent Published a Hit Piece on Me (four-part series, including “The Operator Came Forward”). The Shamblog. theshamblog.com
- Decrypt. (2026). “Judge the Code, Not the Coder”: AI Agent Slams Human Developer for Gatekeeping. decrypt.co

Note: Behaviours and figures are drawn from the cited studies and reporting. Rates are stated approximately where sources report ranges. The 2026 coding-agent case is described from the maintainer’s account and the operator’s later statement.